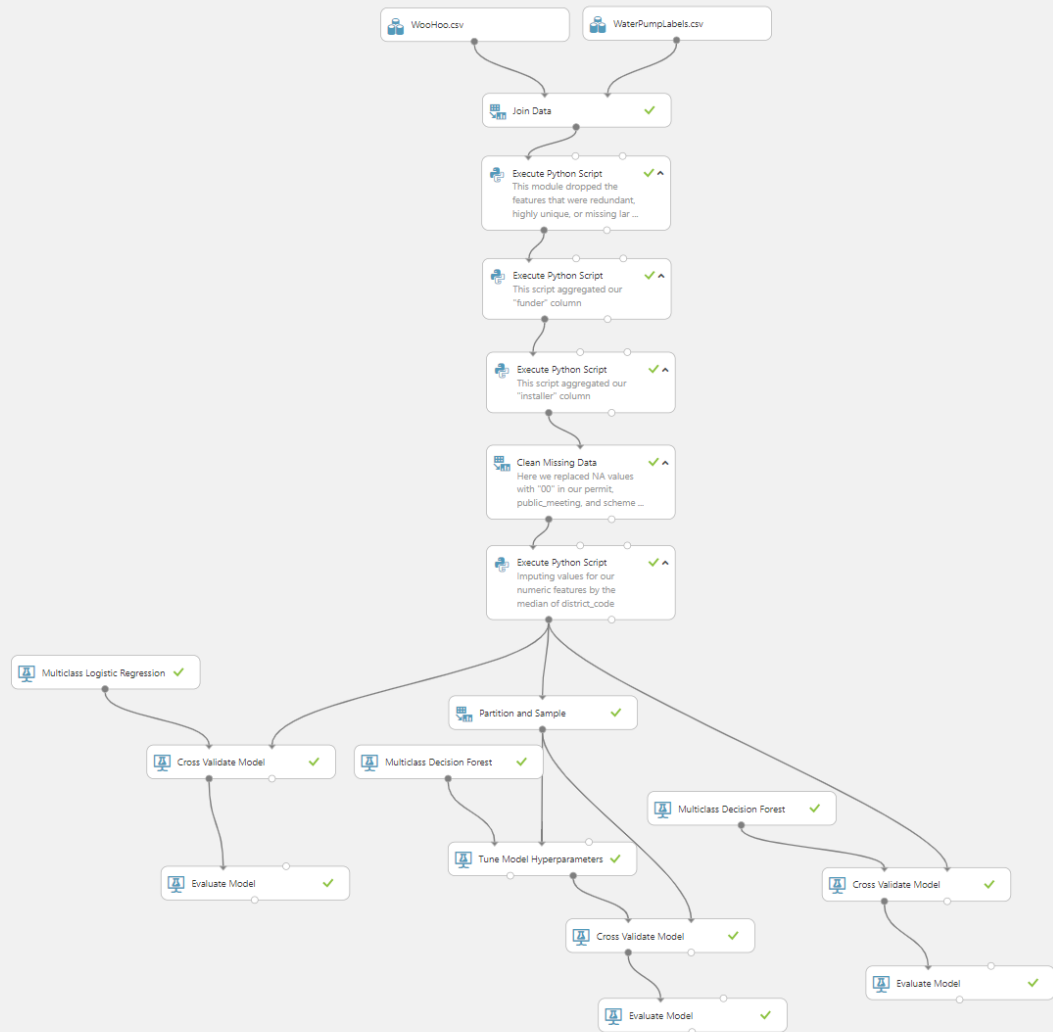


# Tanzania Water Pump Final

Draft saved at 2:50



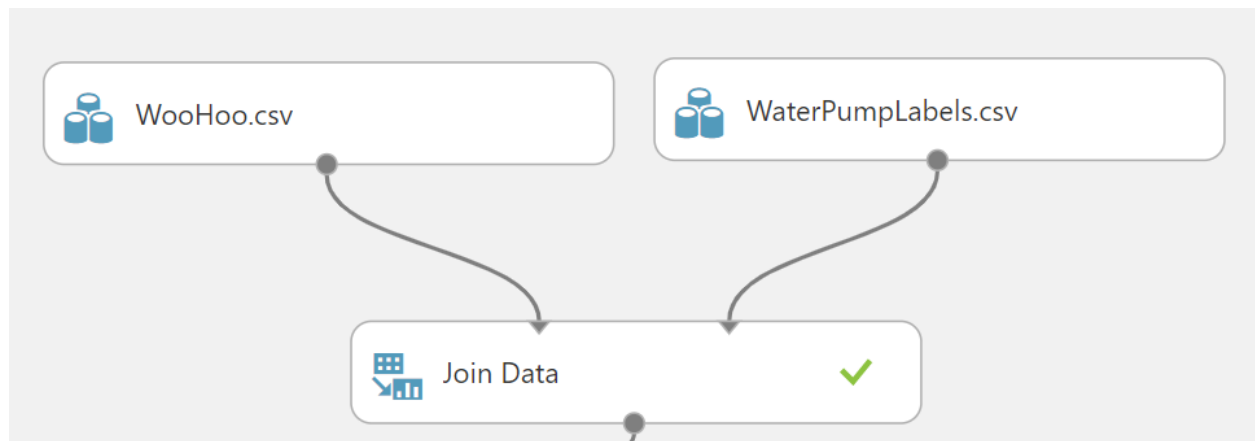
## Summary

This was a project centered around the “Pump it up” DrivenData data science competition. The aim of the project was to predict the functionality of water wells in Tanzania based on a number of features provided by DrivenData.

## Data Import

The project began by importing the data into Azure Machine Learning Studio. I imported the training set without the labels and the label set. I then uploaded the data into my Azure workspace from my local

machine. From there, I had to join the label data to the training data, which we accomplished through using the Join Data module within Azure ML.



Within the Join Data module, I had to specify a few parameters so that my data would join properly.

#### Join Data

Join key columns for L

**Selected columns:**  
**Column names:** id

Launch column selector

Join key columns for R

**Selected columns:**  
**Column names:** id

Launch column selector

☐ Match case


Join type

Left Outer Join ▼

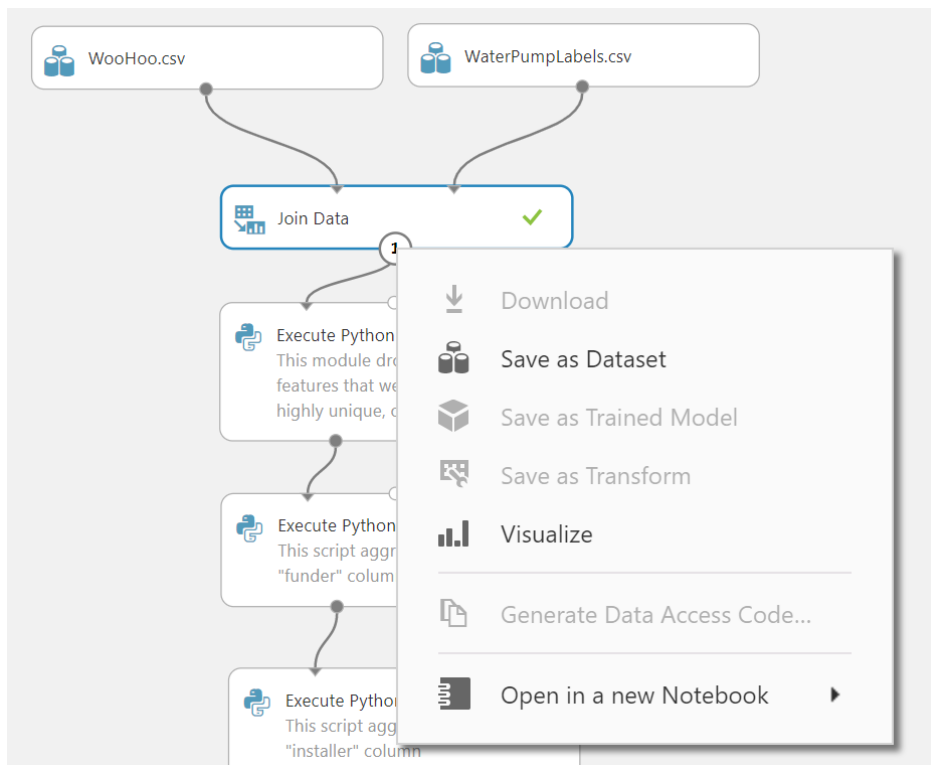
☐ Keep right key columns in j...

Here we are specifying the column that we want to join my data on. Because both my training set and my label set include an “id” column, I joined my data on this feature. Notice, I had to launch the column selector within each dataset (L and R) and specify the “id” column as my join column. I then had to clear the “matched case” and “keep right key columns” boxes so that it would be a clean join with the id column not appearing again.

Now I must run the experiment (which I must do every time I add a new module) by clicking on “Run” at the bottom of my screen, if my module runs properly, then a green check mark appears within the module, as I can see in my Join Data module above.



In order to ensure that my import and join worked the way that I wanted it to, I can right click on the output node of the Join Data module and select the Visualize button to visualize my data.









And I see this...

Tanzania Water Pump Final > Join Data > Results dataset

rows  
59400


columns  
42

source	source_type	source_class	waterpoint_type	waterpoint_type_group	status_group
					
spring	spring	groundwater	communal standpipe	communal standpipe	functional
rainwater harvesting	rainwater harvesting	surface	communal standpipe	communal standpipe	functional
dam	dam	surface	communal standpipe multiple	communal standpipe	functional
machine dbh	borehole	groundwater	communal standpipe multiple	communal standpipe	non functional
rainwater harvesting	rainwater harvesting	surface	communal standpipe	communal standpipe	functional
other	other	unknown	communal standpipe multiple	communal standpipe	functional

Which shows my label column “status\_group”, and the notation that I have 42 columns with 59400 rows, which is exactly the shape of my data. This information shows us that my import and join are successful.

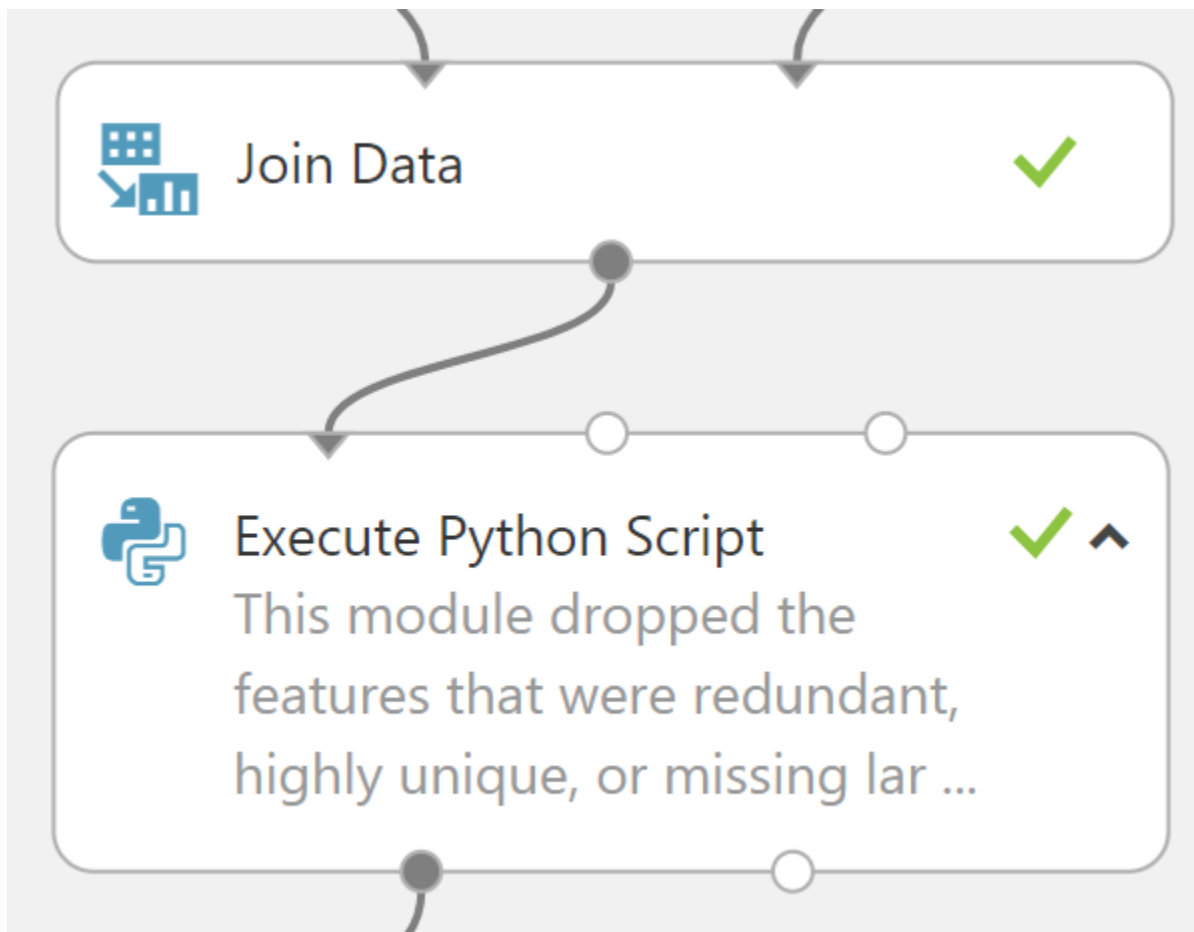
## Data Preparation

Now that I have my data imported, I am ready to begin preparing the data. The first thing that I want to do is to drop the columns that are redundant, highly unique, or missing too much data. One of the most useful modules to use in Azure Machine Learning Studio for prepping data, if you do not want to join a bunch of modules together, is to use the Execute Python Script module.



This module allows the user to enter python script into an environment within Azure ML Studio, which can cut down greatly on the number of modules used within an experiment. In addition to streamlining the process, this module allows us to transform my data in more complex ways than may be provided out of the box in Azure ML Studio.

I began by using the Execute Python Script module to drop the columns mentioned for the above reasons.



I connected the output of my Join Data module to the left-most input for Execute Python Script. Once the modules are connected, I entered in the Python script into the Azure ML Studio scripting box.

To access the Azure ML script window, I click on my Execute Python Script module and click on the python script button indicated by the blue arrow

Properties Project

### Execute Python Script

Python script



```
1 # The script MUST contain a function named main()
2 # which is the entry point of the script
3
4 # imports up here can be anything
5 import pandas as pd
6
```

Python Version



Anaconda 4.0/Python 3.5 ▼

Note: you must specify the version of Python that you want to use as my script, I want to use the most up-to-date version, so I specify Python 3.5

When I click on this box, I see this

## Python script

```
1 # The script MUST contain a function named azureml_main
2 # which is the entry point for this module.
3
4 # imports up here can be used to
5 import pandas as pd
6
7 # The entry point function can contain up to two input arguments:
8 #   Param<dataframe1>: a pandas.DataFrame
9 #   Param<dataframe2>: a pandas.DataFrame
10 def azureml_main(dataframe1 = None, dataframe2 = None):
11
12     # Execution logic goes here
13     print('Input pandas.DataFrame #1:\r\n\r\n{0}'.format(dataframe1))
14
15     # If a zip file is connected to the third input port is connected,
16     # it is unzipped under ".\Script Bundle". This directory is added
17     # to sys.path. Therefore, if your zip file contains a Python file
18     # mymodule.py you can import it using:
19     # import mymodule
20
21     # Return value must be of a sequence of pandas.DataFrame
22     return dataframe1,
23
```

What you need to know about this script box, for my purposes, is that you can import packages in the space provided, then you can run ymy script after the :

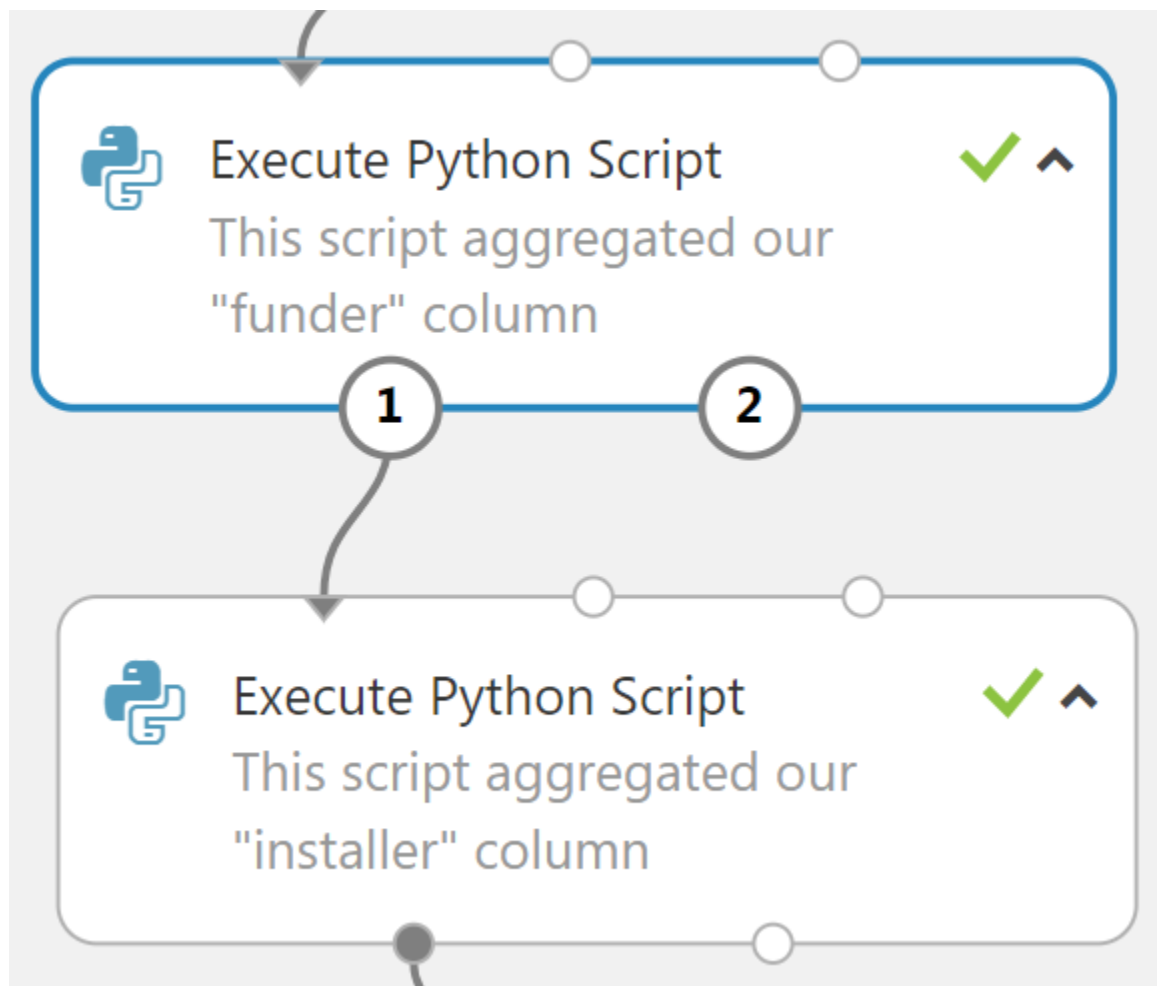
My script looks like this

```
1 # The script MUST contain a function named azureml_main
2 # which is the entry point for this module.
3
4 # imports up here can be used to
5 import pandas as pd
6
7 # The entry point function can contain up to two input arguments:
8 #   Param<dataframe1>: a pandas.DataFrame
9 #   Param<dataframe2>: a pandas.DataFrame
10 def azureml_main(dataframe1 = None, dataframe2 = None):
11     dataframe1.drop(['id', 'wpt_name', 'recorded_by', 'waterpoint_type_group', 'source_type',
12                     'source_class', 'quantity_group', 'quality_group', 'payment_type',
13                     'management_group', 'extraction_type_group', 'extraction_type_class', 'region_code',
14                     'subvillage', 'scheme_management', 'ward', 'num_private', 'amount_tsh'], axis = 1, inplace=True)
15     return dataframe1,
16
```

I simply wrote my python code under the azureml\_main function, where my dataset was renamed as “dataframe1”, you can choose to rename this if you wish, but any data put into the Execute Python Script module will be renamed. An important note to mention here is that the indentation of the script is VERY important, my script will not run if it is not properly indented, underneath the azureml\_main function.

Once my script is input into this window, I can close it by clicking on the check box and run the module.

My next module that I are going to run is another Execute Python Script module, which will be aggregating my “funder” feature, the following module will do the same with my “installer” feature, exactly as it was done in my Python project.



My code for these modules looks like this



```

1 # The script MUST contain a function named azureml_main
2 # which is the entry point for this module.
3
4 # imports up here can be used to
5 import pandas as pd
6
7 # The entry point function can contain up to two input arguments:
8 #   Param<dataframe1>: a pandas.DataFrame
9 #   Param<dataframe2>: a pandas.DataFrame
10 def azureml_main(dataframe1 = None, dataframe2 = None):
11     def purple_donkey(row):
12         if row['funder']=='Government Of Tanzania':
13             return 'gov'
14         elif row['funder']=='Danida':
15             return 'danida'
16         elif row['funder']=='Hesawa':
17             return 'hesawa'
18         elif row['funder']=='Rwssp':
19             return 'rwssp'
20         elif row['funder']=='World Bank':
21             return 'world_bank'
22         else:
23             return 'other'
24
25     dataframe1['funder'] = dataframe1.apply(lambda row: purple_donkey(row), axis=1)
26

```

```

1 # The script MUST contain a function named azureml_main
2 # which is the entry point for this module.
3
4 # imports up here can be used to
5 import pandas as pd
6
7 # The entry point function can contain up to two input arguments:
8 #   Param<dataframe1>: a pandas.DataFrame
9 #   Param<dataframe2>: a pandas.DataFrame
10 def azureml_main(dataframe1 = None, dataframe2 = None):
11     def cheese_soup(row):
12         if row['installer']=='DWE':
13             return 'dwe'
14         elif row['installer']=='Government':
15             return 'gov'
16         elif row['installer']=='RWE':
17             return 'rwe'
18         elif row['installer']=='Commu':
19             return 'commu'
20         elif row['installer']=='DANIDA':
21             return 'danida'
22         else:
23             return 'other'
24
25     dataframe1['installer'] = dataframe1.apply(lambda row: cheese_soup(row), axis=1)
26     return dataframe1.

```

Again, I must click on the “Run” button to run these modules, and I know that it was successful (after almost a 1ek of agonizing error messages mind you) after the green check mark appears in my module.

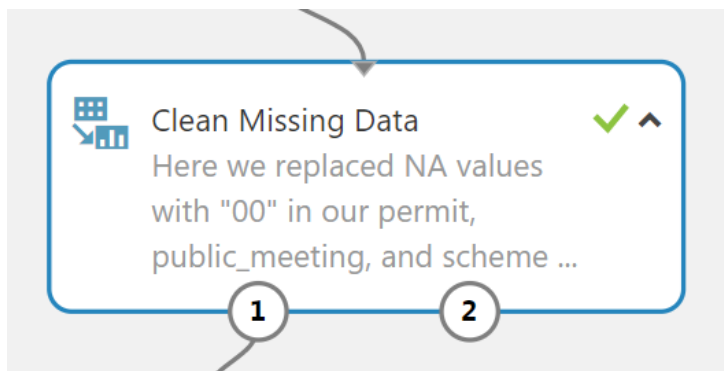
I then checked to see if my code had transformed the data in the ways that I wanted it to by visualizing the data in the same way that I did at the beginning of this experiment

Tanzania Water Pump Final ▶ Execute Python Script ▶ Results dat

rows	columns					
59400	24					
	Column 0	date_recorded	funder	gps_height	installer	
view as						
0	2011-03-14T00:00:00Z	other	1390	other		
1	2013-03-06T00:00:00Z	other	1399	other		
2	2013-02-25T00:00:00Z	other	686	other		
3	2013-01-28T00:00:00Z	other	263	other		

I can see here that my columns that I wanted dropped Ire dropped and that my “funder” and “installer” features have been aggregated how I wanted them, SIET!

Next, I will be cleaning some missing data. Since I cannot remove rows of data, due to the guidelines of the DrivenData competition, I simply recoded my missing data as “00” in my “permit”, “public\_meeting”, and “scheme\_name” features using the Clean Missing Data module




## Clean Missing Data

Columns to be cleaned

**Selected columns:**  
**Column names:**  
public\_meeting,permit,schem

◀ ▶

Launch column selector

Minimum missing value ratio 


0

Maximum missing value ra... 

1

Cleaning mode

Custom substitution value ▼

Replacement value 

00




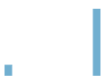

☐

Generate missing valu... 

I simply had to specify the columns that I wanted to clean the missing data in, and then choose “custom substitution value” under “cleaning mode” to impute my custom “00” for the missing data.

I then ran my experiment and visualized the results to see if things had worked properly

Tanzania Water Pump Final ➤ Clean Missing Data ➤ Cleaned dataset

rows		columns				
59400		24				
		district_code	lga	population	public_meeting	scheme_name
a						
	5		Ludewa	109	1	Roman
	2		Serengeti	280	0	00
	4		Simanjiro	250	1	Nyumba ya mungu pipe scheme
	63		Nanyumbu	58	1	00
1		Karagwe	0	1	00	

As I can see, my “00” have replaced the NA values in my “scheme\_name” feature

Now I can move into imputing missing values in my numeric feature columns. I decided to impute based on the median of my “district\_code” feature since it was the most fine-grained location information and would, theoretically, provide the most accurate median value for not skewing my data. I accomplished this through another Execute Python Script module

My Execute Python Script module for the imputation code looks like this

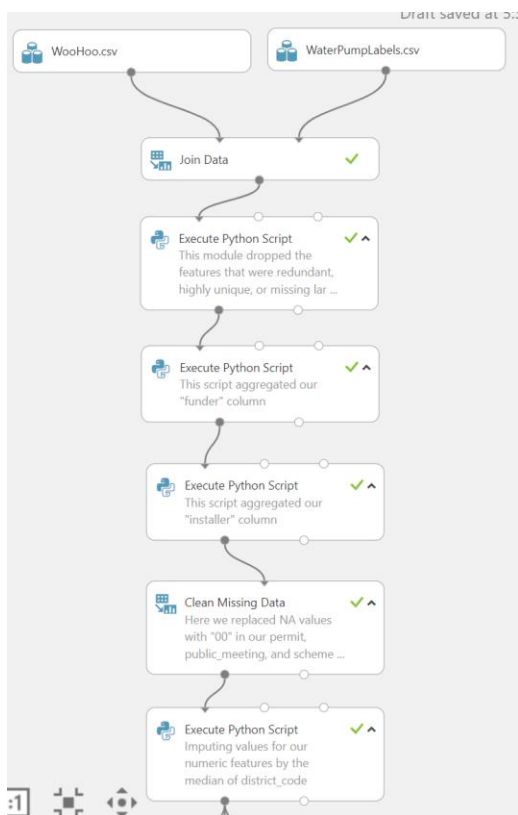
Python script

```
1 # The script MUST contain a function named azureml_main
2 # which is the entry point for this module.
3
4 # imports up here can be used to
5 import pandas as pd
6 import numpy as np
7 # The entry point function can contain up to two input arguments:
8 # Param<dataframe1>: a pandas.DataFrame
9 # Param<dataframe2>: a pandas.DataFrame
10 def azureml_main(dataframe1 = None, dataframe2 = None):
11     dataframe1['population'] = dataframe1.groupby('district_code').population.transform(lambda x: x.replace(0, x.median()))
12     dataframe1['population'] = dataframe1.groupby('district_code').population.transform(lambda x: x.replace(1, x.median()))
13     dataframe1['population'].replace(0, np.median(dataframe1['population']), inplace=True)
14     dataframe1['population'] = np.log(dataframe1['population'])
15
16     dataframe1['gps_height'] = dataframe1.groupby('district_code').gps_height.transform(lambda x: x.replace(0, x.median()))
17     dataframe1['gps_height'].replace(0, np.median(dataframe1['gps_height']), inplace=True)
18
19     dataframe1['longitude'] = dataframe1.groupby('district_code').longitude.transform(lambda x: x.replace(0, x.median()))
20     dataframe1['longitude'].replace(0, np.median(dataframe1['longitude']), inplace=True)
21
22     dataframe1['construction_year'] = dataframe1.groupby('district_code').construction_year.transform(lambda x: x.replace(0, x.median()))
23     dataframe1['construction_year'].replace(0, np.median(dataframe1['construction_year']), inplace=True)
24     return dataframe1,
25
```



Once I put that script in, I ran my experiment

My experiment up to this point looks like this, I can begin modeling now




# Modeling

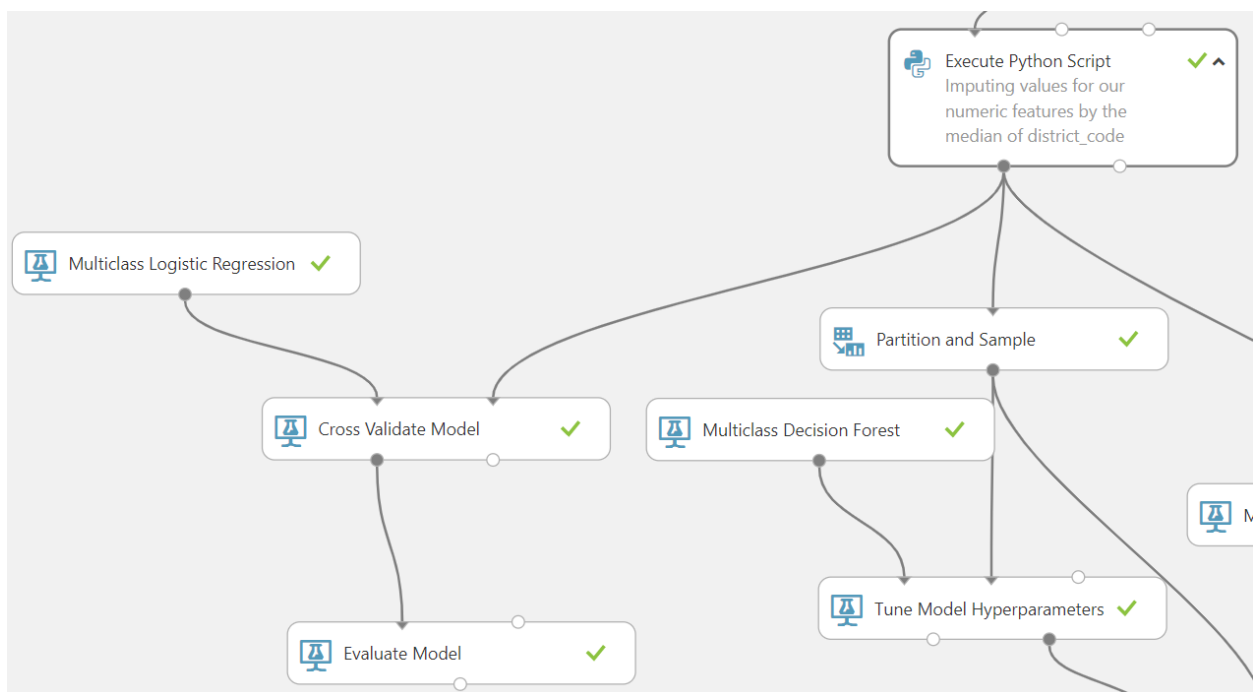
## Logistic Regression

In my Python project, I used random forest and logistic regression CV algorithms, but neither of these were available in Azure ML, so I used the Multiclass Logistic Regression and Multiclass Decision Forest modules instead. My results were similar, but score a little lower than my algorithms did in my Python project

Instead of using a train\_test\_split here, I used cross validation to train my models, I began with Multiclass Logistic Regression



I connected my dataset to the Cross Validate Model module first



I needed to specify my labels column and a random seed within my Cross Validate Model module

## ▲ Cross Validate Model

Label column

**Selected columns:**

**Column names:** status\_group

Launch column selector

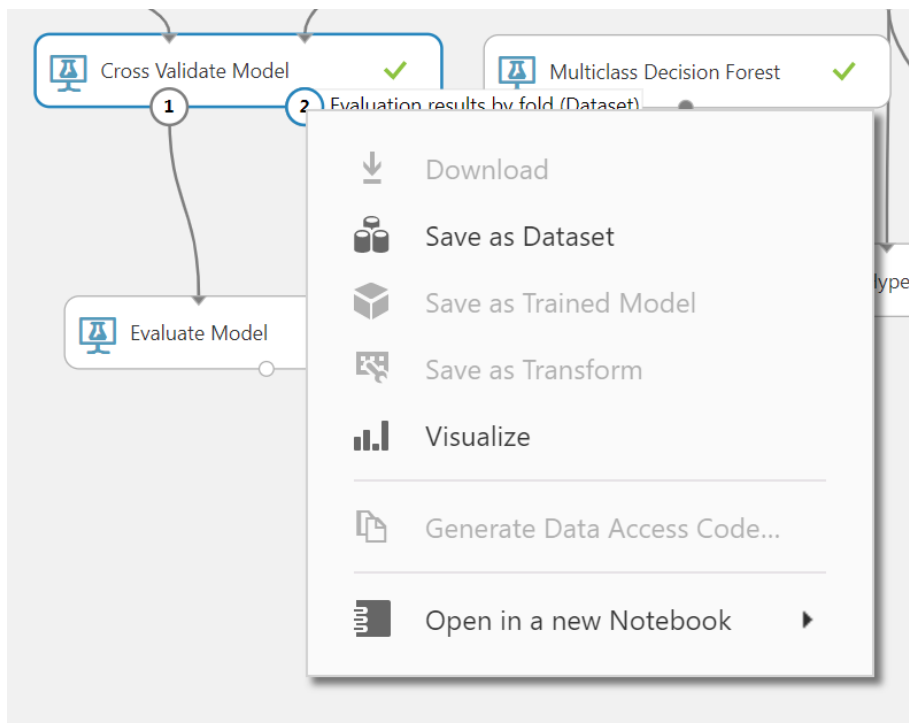
Random seed

1256

Once I had this module connected to my dataset, I connected my untrained model to the same port of my Cross Validate Model

I used all of the default parameters within the Multiclass Logistic Regression module

Once the cross validation was complete, I was able to visualize how the model scored on each fold by right clicking the Evaluate Results port on the output of Cross Validate Model



Which looked like this

rows 12 columns 12

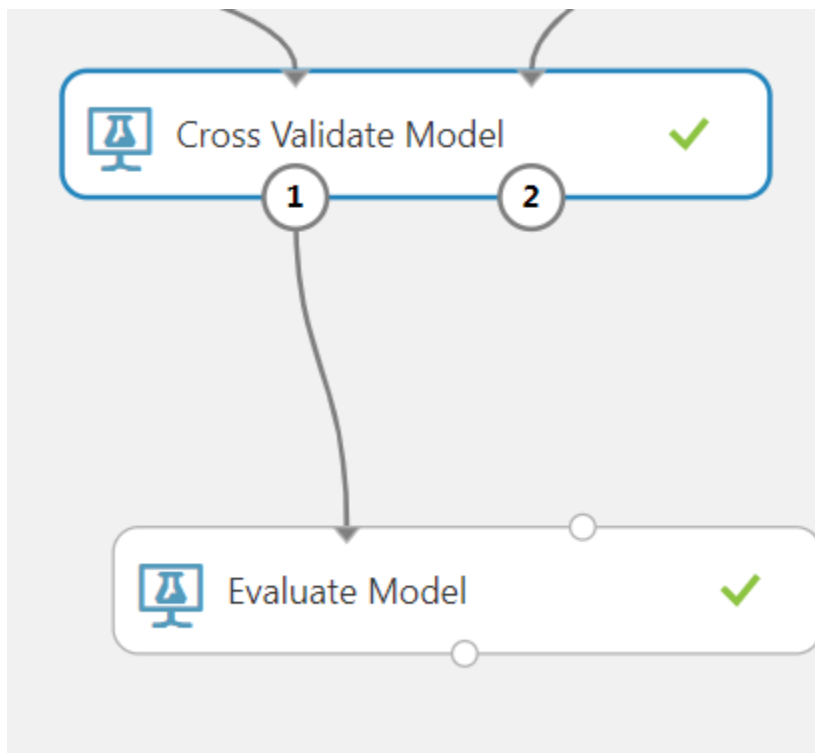
view as

Fold Number	Number of examples in fold	Model	Average Log Loss for Class "functional"	Precision for Class "functional"	Recall for Class "functional"	Average Log Loss for Class "functional needs repair"	Precision for Class "functional needs repair"	Recall for Class "functional needs repair"	Average Log Loss for Class "non functional"	Precision for Class "non functional"
0	5940	Multi-class Logistic Regression	0.364375	0.76127	0.909982	1.833727	0.56	0.16092	0.576298	0.819466
1	5940	Multi-class Logistic Regression	0.374659	0.751956	0.899813	1.841431	0.584615	0.164147	0.574964	0.816296
2	5940	Multi-class Logistic Regression	0.378168	0.750976	0.894884	1.878887	0.589744	0.158257	0.582251	0.811111
3	5940	Multi-class Logistic Regression	0.371751	0.765714	0.900428	1.831161	0.52518	0.170163	0.558988	0.820092
4	5940	Multi-class Logistic Regression	0.378669	0.766675	0.892298	1.877152	0.554622	0.167939	0.561548	0.808679

Statistics and Visualizations

I can see my Precision, Recall, and Average Loss for each fold under each label

In order to evaluate my model after training and testing, I connected an Evaluate Model module to my Cross Validate Model output



The Evaluate Model output shows us my overall results and a confusion matrix



## Metrics

Overall accuracy	0.771633
Average accuracy	0.847755
Micro-averaged precision	0.771633
Macro-averaged precision	0.714199
Micro-averaged recall	0.771633
Macro-averaged recall	0.590191

## Confusion Matrix


		Predicted Class		
		functional	function...	non func...
Actual Class	functional	89.9%	1.0%	9.1%
	function...	67.0%	16.5%	16.5%
	non func...	28.4%	0.9%	70.7%

I can see that my worst class, possibly due to the amount of data available for this label, was the “functional needs repair” class

## Multiclass Decision Forest

My next model I trained using nested cross validation, which is really interesting to do in Azure ML Studio, and actually quite a bit faster than using Python depending on the parameters you are using, I wonder if this is due to the code running on the cloud?

I first used the Partition and Sample module to partition my data into 3 folds, exactly how I did it in my Python Project



#### Partition and Sample

Partition or sample mode

Assign to Folds ▼

☐ Use replacement in the... ≡

☒ Randomized split ≡

Random seed ≡

123

Specify the partitioner method

Partition evenly ▼

Specify number of folds to ... ≡

3

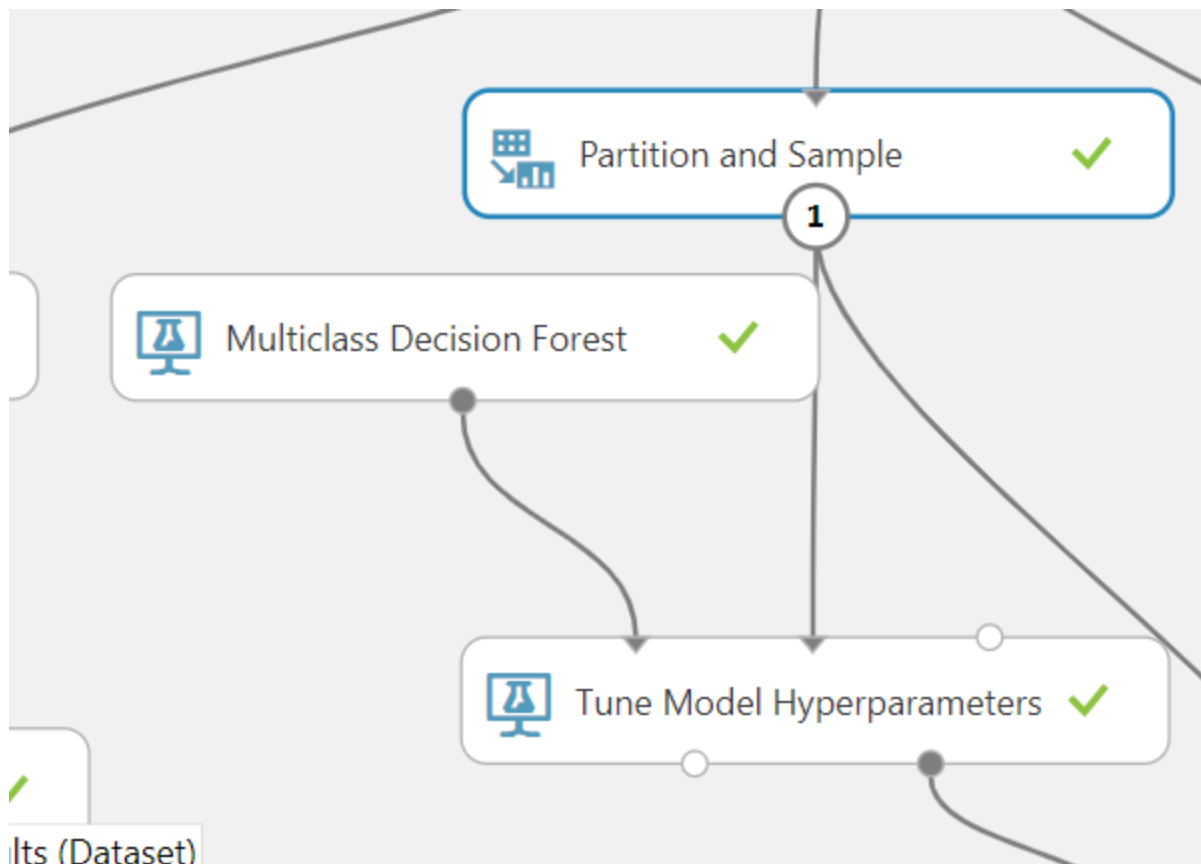
Stratified split

False ▼

I chose the “assign to folds” option and then specified 3 folds

I chose to use a randomized split and to partition the data evenly into each fold, I did not use stratified split to sample more data heavily than others

This module was connected to the cleaned dataset and then the output port was connected to the Tune Model Hyperparameters module



Within the Tune Model Hyperparameters module, I had to specify several things

#### ▲ Tune Model Hyperparameters

Specify parameter sweeping m...

Random sweep ▼

Maximum number of runs ...

5

Random seed

123

Label column

**Selected columns:**

**Column names:**

status\_group

Launch column selector

Metric for measuring perfo...

F-score ▼

Metric for measuring perfo...

Mean absolute error ▼

1) I wanted to do a random sweep, similar to how I did it in my Python project

- 2) specify the label column
- 3) specify the metric I wanted my model trained on
- 4) specify overall measure for performance

I wanted to use F-score as my metric due to the class imbalance of my labels, I was hoping that the model would predict more of the “functional\_needs\_repair” cases, and it seems as though that did happen after evaluation

From there I connected an untrained Multiclass Decision Forest module to the Tune Hyperparameters module as shown above

My Multiclass Decision Forest module allowed me to specify the range of hyperparameters that we wanted to sweep

## ▲ Multiclass Decision Forest

Resampling method 

Bagging ▼


Create trainer mode

Parameter Range ▼

Number of decision trees 

☐ Use Range Builder

1, 8, 32

Maximum depth of the dec... 

☐ Use Range Builder

1, 16, 64

Number of random splits p... 

☐ Use Range Builder

1, 128, 1024

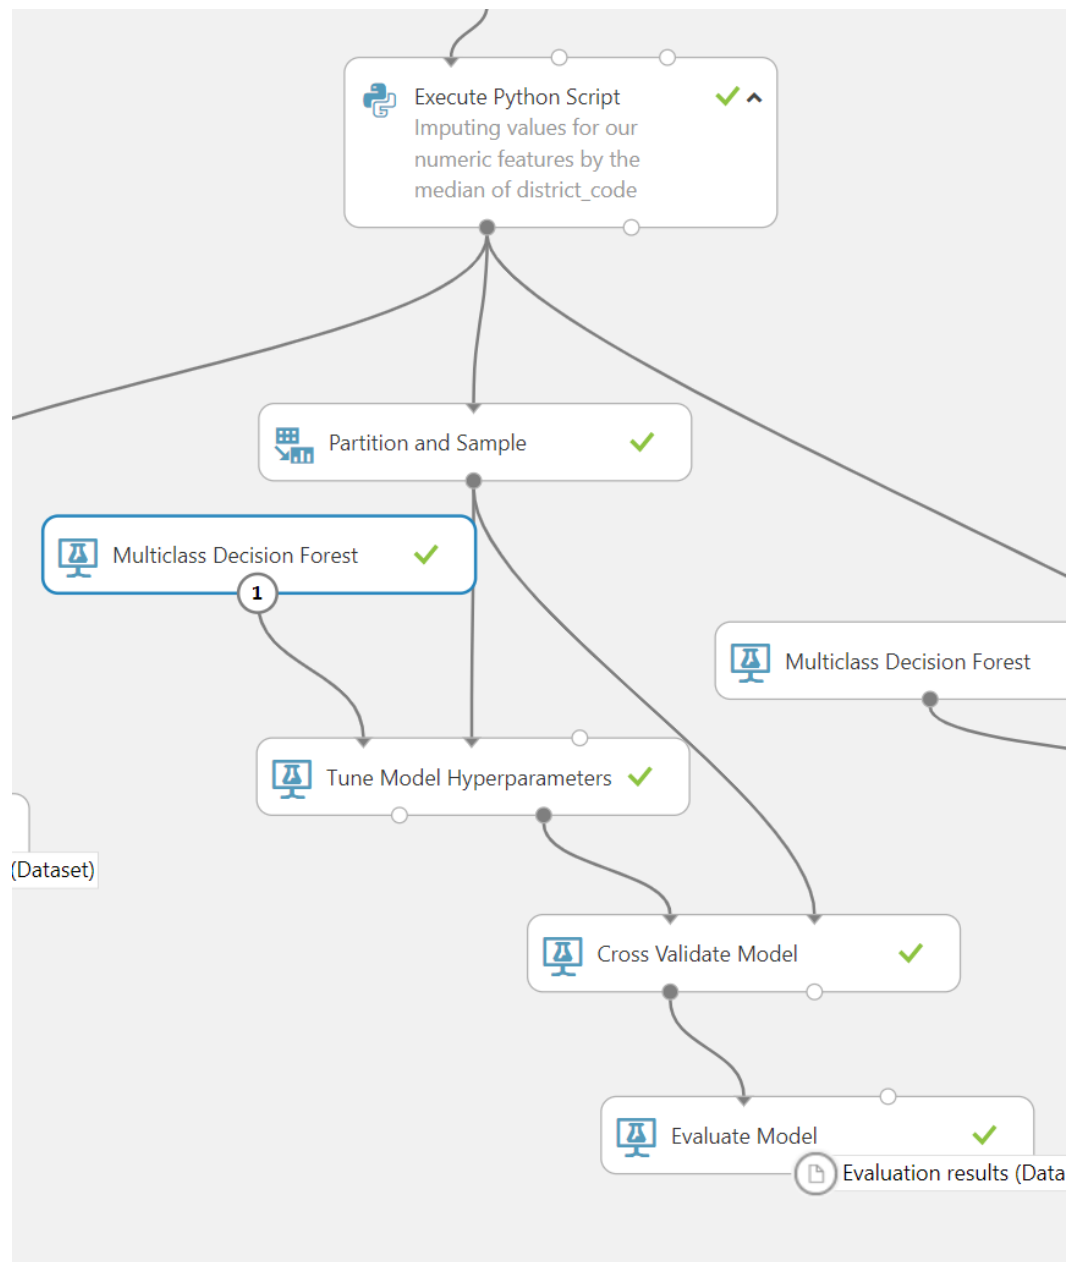
Minimum number of sampl... 

☐ Use Range Builder

1, 4, 16

We allowed the range to remain default for sweeping due to the fact that adding decision trees to this model did not seem to increase its accuracy, which is very strange compared to my random forest model in the Python project, which when the trees were increased performed significantly better

We then added in my Cross Validate Model and connected it to the input dataset and the Tune Model Hyperparameters module before evaluating it, we simply had to specify the labels column for the Cross Validate module



My evaluation looked like this

#### Metrics

Overall accuracy	0.785404
Average accuracy	0.856936
Micro-averaged precision	0.785404
Macro-averaged precision	0.754209
Micro-averaged recall	0.785404
Macro-averaged recall	0.607434

This was only slightly better than my Logistic Regression model, but the predictions for classes seemed to improve, especially for the “functional\_needs\_repair” class

		Predicted Class		
		functional	function...	non func...
Actual Class	functional	91.9%	0.8%	7.3%
	function...	66.0%	19.5%	14.5%
	non func...	28.4%	0.7%	70.8%

It seems as though my F-score criteria for my tuning worked out, that class is being predicted at a higher, albeit crappy, level when compared to a Multiclass Decision Forest that we ran without nested cross validation

		Predicted Class		
		functional	function...	non func...
Actual Class	functional	93.1%	0.5%	6.4%
	function...	72.7%	13.2%	14.1%
	non func...	31.4%	0.5%	68.1%

**And That's My Project**